



第12章 交换网络中的路由选择

南京大学计算机系 黄皓教授

2007年9月30日 星期五



12.1 Routing in Circuit Switched Network

- Many connections will need paths through more than one switch
- Need to find a route
 - Efficiency
 - Resilience
- Static routing
 - Public telephone switches are a tree structure
 - Static routing uses the same approach all the time
- Dynamic routing allows for changes in routing depending on traffic
 - Uses a peer structure for nodes




Alternate Routing

- Possible routes between end offices predefined
- Originating switch selects appropriate route
- Routes listed in preference order
- Different sets of routes may be used at different times



Alternate Routing Diagram



Route a: X® Y
Route b: X® J® Y
Route c: X® K® Y
Route d: X® I® J® Y

○ = end office

⬡ = intermediate switching node

(a) Topology

Time Period	First route	Second route	Third route	Fourth and final route
Morning	a	b	c	d
Afternoon	a	d	b	c
Evening	a	d	c	b
Weekend	a	c	b	d

(b) Routing table




12.2 Routing in Packet Switched Network

- Complex, crucial aspect of packet switched networks
- Characteristics required
 - Correctness
 - Simplicity
 - Robustness
 - Stability
 - Fairness
 - Optimality
 - Efficiency



Example Packet Switched Network






12.2.1 Elements of Routing

- Performance criteria
- Decision time
- Decision place
- Network info source
- Network info update timing



Performance Criteria

- Minimum hop
 - e.g. 1–3–6
- Least cost
 - e.g. 1–4–5–6
- Others
 - Minimum delay
 - Largest throughput





Decision Time and Place

■ Time

- ☐ Packet
- ☐ virtual circuit

■ Place

- ☐ Distributed
 - Made by each node
- ☐ Centralized
- ☐ Source



Network Information Source and Update Timing

- Routing decisions usually based on knowledge of network (not always)
- Distributed routing
 - Nodes use local knowledge
 - May collect info from adjacent nodes
 - May collect info from all nodes on a potential route
- Central routing
 - Collect info from all nodes
- Update timing
 - When is network info held by nodes updated
 - Fixed - never updated
 - Adaptive - regular updates



Conclusion of Routing Elements

Performance Criteria

- Number of hops
- Cost
- Delay
- Throughput

Network Information Source

- None
- Local
- Adjacent node
- Nodes along route
- All nodes

Decision Time

- Packet (datagram)
- Session (virtual circuit)

Network Information Update Timing

- Continuous
- Periodic
- Major load change
- Topology change

Decision Place

- Each node (distributed)
- Central node (centralized)
- Originating node (source)



12.2.2 Routing Strategies

- Fixed
- Flooding
- Random
- Adaptive



(1) Fixed Routing

- Single permanent route for each source to destination pair
- Determine routes using a least cost algorithm
- Route fixed, at least until a change in network topology



(1) Fixed Routing tables

CENTRAL ROUTING DIRECTORY

From Node

	1	2	3	4	5	6
1	—	1	5	2	4	5
2	2	—	5	2	4	5
3	4	3	—	5	3	5
4	4	4	5	—	4	5
5	4	4	5	5	—	5
6	4	4	5	5	6	—

To Node

Node 1 Directory

Destination	Next Node
2	2
3	4
4	4
5	4
6	4

Node 2 Directory

Destination	Next Node
1	1
3	3
4	4
5	4
6	4

Node 3 Directory

Destination	Next Node
1	5
2	5
4	5
5	5
6	5

Node 4 Directory

Destination	Next Node
1	2
2	2
3	5
5	5
6	5

Node 5 Directory

Destination	Next Node
1	4
2	4
3	3
4	4
6	6

Node 6 Directory

Destination	Next Node
1	5
2	5
3	5
4	5
5	5



(2) Flooding

- No network info required
- Packet sent by node to every neighbor
 - Incoming packets retransmitted on every link except incoming link
- Eventually a number of copies will arrive at destination
- Each packet is uniquely numbered so duplicates can be discarded
- Nodes can remember packets already forwarded to keep network load in bounds
- Can include a hop count in packets



(2) Flooding Example

Hop count = 3

- Initial


- 3 packets

- 1st hop


- 9 packets

- 2nd hop


- 23 packets



(a) First hop



(b) Second hop



(c) Third hop



(2) Properties of Flooding

- All possible routes are tried
 - Very robust
- At least one packet will have taken minimum hop count route
 - Can be used to set up virtual circuit
- All nodes are visited
 - Useful to distribute information (e.g. routing)



(3) Random Routing

- Node selects one outgoing path for retransmission of incoming packet
- Selection can be random or round robin
- Can select outgoing path based on probability calculation
- No network info needed
- Route is typically not least cost nor minimum hop



(3) Random Routing

Assign Probabilities

- $P_i = R_i / \sum_j R_j$
 - P_i – Probability of selecting out-link i
 - R_i – Cost factor of link i
- Possible **cost factor**
 - Transmission rate – for throughput
 - Queue size – for delay
- Partly used in **adaptive routing**



Adaptive Routing

- Used by almost all packet switching networks
- Routing decisions change as conditions on the network change
 - Failure
 - Congestion
- Requires info about network
- Decisions more complex
- Tradeoff between quality of network info and overhead
- Reacting too quickly can cause oscillation
- Too slowly to be relevant



Adaptive Routing - Advantages

- Improved performance
- Aid congestion control
- Complex system
 - May not realize theoretical benefits



Classification


- Based on information sources
 - Local (isolated)
 - Route to outgoing link with shortest queue
 - Can include bias for each destination
 - Rarely used - do not make use of easily available info
 - Adjacent nodes
 - All nodes



Isolated Adaptive Routing

Node 4's Bias
Table for
Destination 6

Next Node	Bias
1	9
2	6
3	3
5	0





12.3 Least Cost Algorithms

- Given network of nodes connected by bi-directional links
- Each link has a cost in each direction
 - Each link cost 1 – **minimum hop count**
 - Link cost inversely proportional to capacity – **maximum throughput**
- Link costs in different directions may be different
 - e.g. length of packet queue
- Define **cost of path** between two nodes as sum of costs of links traversed
- For each pair of nodes, find a path with the least cost



2 Algorithms in Adaptive Routing

- Dijkstra's Algorithm
- Bellman-Ford Algorithm



Dijkstra's Algorithm

- Find shortest paths **from given source to all other nodes**
 - Developing paths in order of increasing path cost (length)
- $N =$ set of nodes in the network
- $s =$ source node
- $T =$ set of nodes so far incorporated by the algorithm
- $w(i, j) =$ link cost from node i to node j
 - $w(i, i) = 0$
 - $w(i, j) = \infty$ if the two nodes are not directly connected
 - $w(i, j) \geq 0$ if the two nodes are directly connected
- $L(n) =$ **cost of least-cost path** from node s to node n currently known
 - At termination, $L(n)$ is cost of least-cost path from s to n



Dijkstra's Algorithm Method

■ Step 1 [Initialization]

- $T = \{s\}$ Set of nodes so far incorporated consists of only source node
- $L(n) = w(s, n)$ for $n \neq s$
- Initial path costs to neighboring nodes are simply link costs



■ Step 2 [Get Next Node]

- Find node x not in T with least-cost path from s (i.e. $\min L(x)$)
- Incorporate node x into T
- Also incorporate the edge that links x with the node in T that contributes to the path

■ Step 3 [Update Least-Cost Paths]

- $L(n) = \min[L(n), L(x) + w(x, n)]$ for all $n \notin T$
- If latter term is minimum, path from s to n is path from s to x concatenated with link from x to n

- Algorithm terminates when
all nodes have been added to T











Dijkstra's Algorithm Notes

- One iteration of steps 2 and 3 adds one new node to T
 - Defines least cost path from s to that node
- At termination, value $L(n)$ associated with each node n is the cost (length) of least-cost path from s to n
- In addition, T defines the least-cost path from s to each other node




Example of Dijkstra's Algorithm





Results of Example Dijkstra's Algorithm

No	T	L(2)	Path	L(3)	Path	L(4)	Path	L(5)	Path	L(6)	Path
1	{1}	2	1-2	5	1-3	1	1-4	∞	—	∞	—
2	{1,4}	2	1-2	4	1-4-3	1	1-4	2	1-4-5	∞	—
3	{1, 2, 4}	2	1-2	4	1-4-3	1	1-4	2	1-4-5	∞	—
4	{1, 2, 4, 5}	2	1-2	3	1-4-5-3	1	1-4	2	1-4-5	4	1-4-5-6
5	{1, 2, 3, 4, 5}	2	1-2	3	1-4-5-3	1	1-4	2	1-4-5	4	1-4-5-6
6	{1, 2, 3, 4, 5, 6}	2	1-2	3	1-4-5-3	1	1-4	2	1-4-5	4	1-4-5-6



Destination	Next-Hop	Distance
2	2	2
3	4	3
4	4	1
5	4	2
6	4	4



Bellman-Ford Algorithm Definitions

- Find shortest paths from given node containing **at most 1 link**
- Find the shortest paths that containing **at most 2 links**, based on the result of 1 link
- Find the shortest paths of **3 links** based on result of 2 links, and so on
- s = source node
- $w(i, j)$ = link cost from node i to node j
 - $w(i, i) = 0$
 - $w(i, j) = \infty$ if the two nodes are not directly connected
 - $w(i, j) \geq 0$ if the two nodes are directly connected
- h = **maximum number of links** in path at current stage of the algorithm
- $L_h(n)$ = cost of least-cost path from s to n under constraint of no more than h links



Bellman-Ford Algorithm Method

- Step 1 [Initialization]
 - $L_0(n) = \infty$, for all $n \neq s$
 - $L_1(n) = w(s, n)$
 - $L_h(s) = 0$, for all h
- Step 2 [Update]
 - For each successive $h \geq 0$
 - For each $n \neq s$, compute $L_{h+1}(n) = \min_j [L_h(j) + w(j, n)]$
 - Connect n with predecessor node j that achieves minimum
 - Eliminate any connection of n with different predecessor formed during earlier iterations
 - Path from s to n terminates with link from j to n
- Repeat until no change made to route (convergence)







Bellman-Ford Algorithm Notes

- For each iteration with $h=K$ and for each destination node n
 - Compares newly computed paths from s to n of length K with path from previous iteration
- If previous path shorter it is retained
- Otherwise new path is defined




Example of Bellman-Ford Algorithm





Results of Bellman-Ford Example

h	$L_h(2)$	Path	$L_h(3)$	Path	$L_h(4)$	Path	$L_h(5)$	Path	$L_h(6)$	Path
0	∞	—	∞	—	∞	—	∞	—	∞	—
1	2	1-2	5	1-3	1	1-4	∞	—	∞	—
2	2	1-2	4	1-4-3	1	1-4	2	1-4-5	10	1-3-6
3	2	1-2	3	1-4-5-3	1	1-4	2	1-4-5	4	1-4-5-6
4	2	1-2	3	1-4-5-3	1	1-4	2	1-4-5	4	1-4-5-6






Comparison of Info Gathering

■ Bellman-Ford

- Calculation for node n involves knowledge of **link cost to all neighbours** plus **total cost to each neighbour from s**



- Each node can maintain set of **costs and paths** to every other node
- Can exchange info with direct neighbours
- Update **costs and paths** based on info from neighbours and knowledge of direct link costs

■ Dijkstra

- Each node needs **complete topology**
- Must know link costs of all links in network
- Must exchange information with all other nodes



Comparison of Efficiency

■ Evaluation

- Dependent on **processing time** of algorithms
- Dependent on **amount of info** required from other nodes

■ Implementation specific

- Both converge to same solution under **static topology and costs**
- **If link costs change**, algorithms will attempt to catch up
- If link costs depend on traffic, which depends on routes chosen, then feedback
- Both will result in **instability**



12.4 ARPANET Routing Strategies

- First Generation, 1969
- Second Generation, 1979
- Third Generation, 1987



First Generation

■ Method

- Distributed adaptive, uses **Bellman-Ford** algorithm
- Estimated delay using performance criterion **per 128ms**
- Node exchanges delay vector with neighbors
- Update routing table based on incoming info

■ Problem

- Doesn't consider line speed, **just queue length** (lines speed is not similar now)
- Queue length not a good measurement of delay
- Responds **slowly to congestion**



The 1st Routing Algorithm

Desti- nation	Delay	Next node
1	0	—
2	2	2
3	5	3
4	1	4
5	6	3
6	8	3

D_1 S_1

(a) Node 1's Routing
table before update

3
0
3
2
3
5

D_2

7
4
0
2
1
3

D_3

5
2
2
0
1
3

D_4

(b) Delay vectors sent to node 1
from neighbor nodes

Desti- nation	Delay	Next node
1	0	—
2	2	2
3	3	4
4	1	4
5	2	4
6	4	4

$$I_{1,2} = 2$$

$$I_{1,3} = 5$$

$$I_{1,4} = 1$$

(c) Node 1's routing table
after update and link
costs used in update



Second Generation

■ Method


- ☐ Uses **measured delay** as performance criterion
- ☐ Per 10s, the node computes the average delay on each outgoing link
 - Time of retransmit - Time of arrive + Transmission time + Transport time
- ☐ Uses Dijkstra's algorithm
- ☐ The delay info of each link is **sent to all other nodes using flooding**

■ Problem

- ☐ Good under light and medium loads
- ☐ Under **heavy loads**, little correlation between reported delays and those experienced



The Vibration Problem





Third Generation

■ Goal

- Let **some stay on loaded link** to balance and avoid oscillations
- Under heavy load, should give the majority route a good path instead of every route the best path

■ Method

- **Link cost calculations** changed
- Leveling based on current value and previous results
- Use **hop normalized metric** to calculate cost

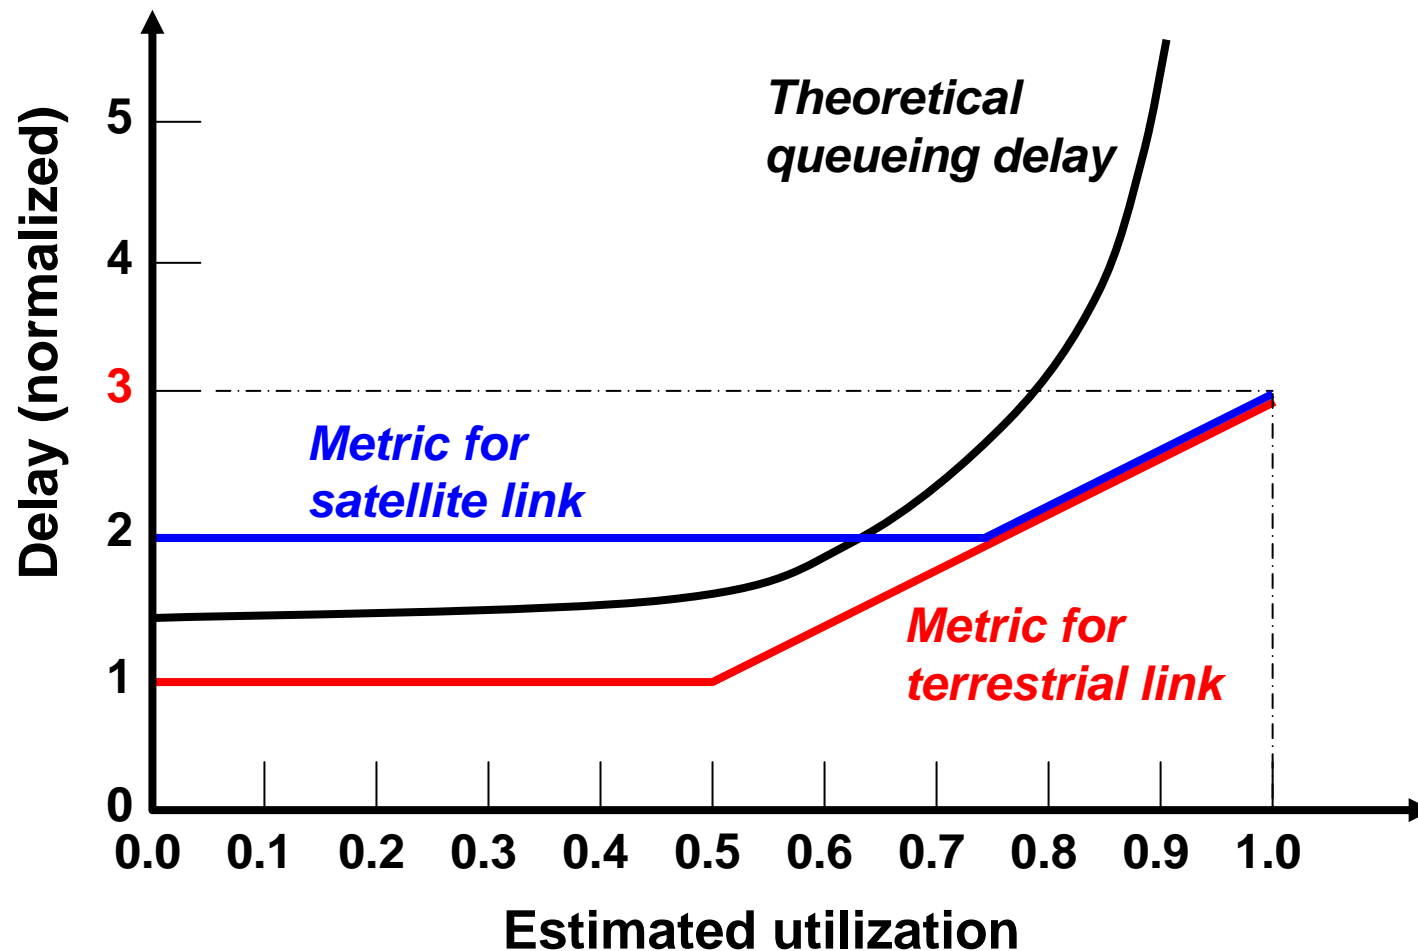


Cost Calculation

- Uses the single-server queuing model, computes **link utilization** instead of delay
- Step 1. **link utilization**
 - $\rho = \frac{2(T_s - T)}{(T_s - 2T)}$
 - ρ – link utilization
 - T – current measured delay
 - T_s – mean packet length (600 bit) / transmission rate of the link
- Step 2. **Leveling**
 - $U_{n+1} = \alpha \times \rho_{n+1} + (1 - \alpha) \times U_n$
 - U_n – average link utilization at time n
 - α – constant, now set 0.5
- Step 3. Set link cost based on **leveled utilization**

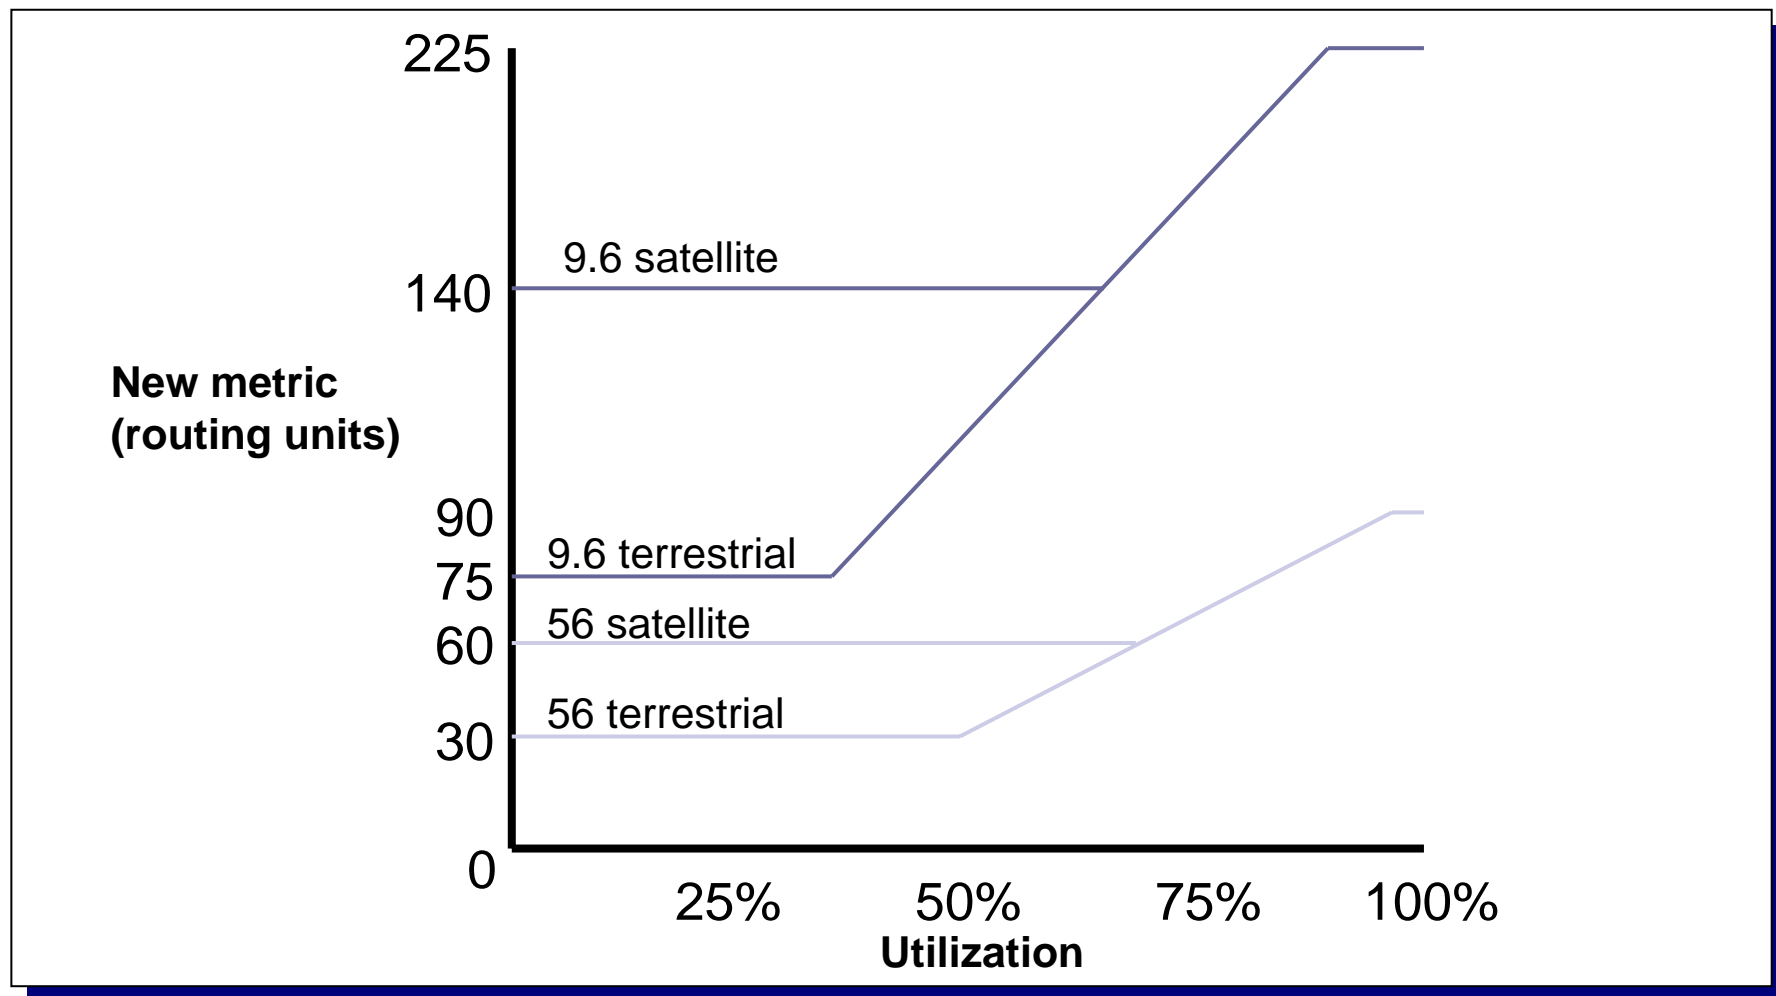


Link Utilization to Cost (1)





Link Utilization to Cost (2)






12.5 Routing Information Protocol



Convergence

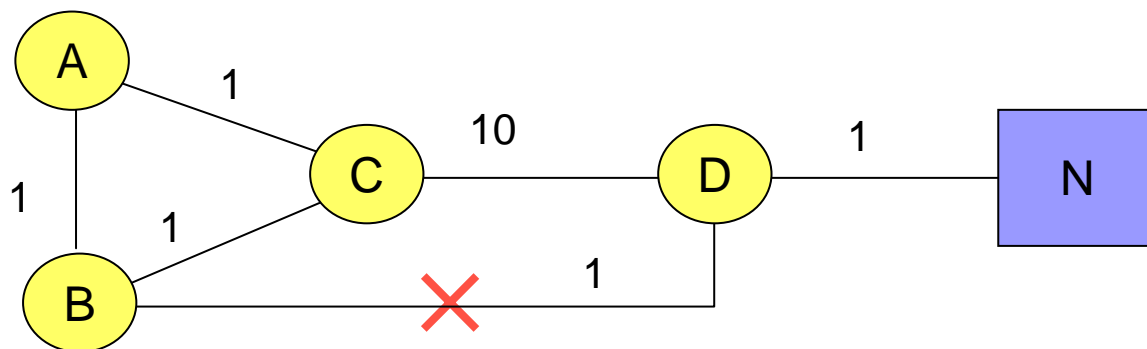




- Router D to the target network: Directly connected network. Metric 1.
- Router B to the target network: Next hop is router D. Metric is 2.
- Router C to the target network: Next hop is router B. Metric is 3.
- Router A to the target network: Next hop is router B. Metric is 3.



counting to infinity

- the link connecting router B and router D fails.

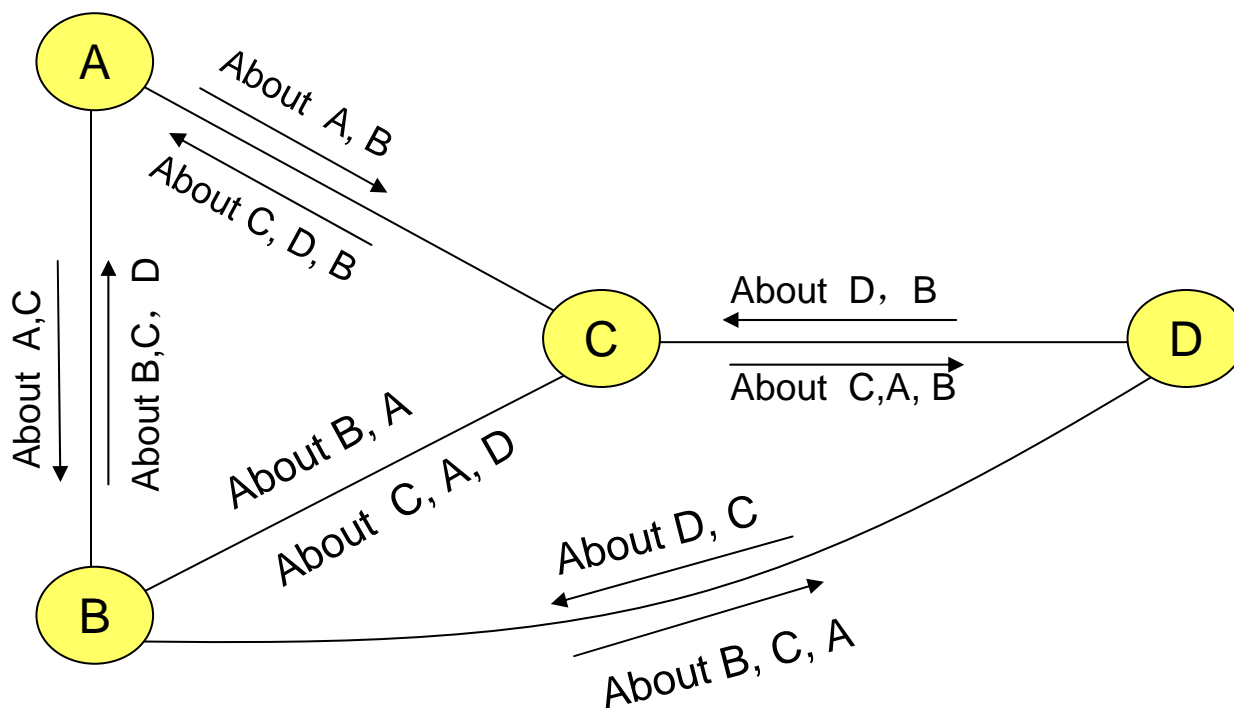


Time																
D: Direct	1	Direct	1	Direct	1	Direct	1	Direct	1	Direct	1				
B: Unreachable		C	4	C	5	C	6		C	11	C	12				
C: B	3	A	4	A	5	A	6		A	11	D	11				
A: B	3	C	4	C	5	C	6	C	11	C	12				




split horizon

- The "simple" scheme omits routes learned from one neighbor in updates sent to that neighbor.





split horizon



Time → →

D: Direct	1	Direct	1	Direct	1	Direct	1
B: Unreachable		Unreachable		Unreachable		C	12
C: B	3	A	4	D	11	D	11
A: B	3	C	4	Unreachable		C	12


Note: Faster Routing Table Convergence

Wait for
timeout



Split horizon with poisoned reverse

- "Split horizon with poisoned reverse" includes such routes in updates, but sets their metrics to infinity.
- If A thinks it can get to D via C, its messages to C should indicate that D is unreachable.
- If the route through C is real, then C either has a direct connection to D, or a connection through some other gateway.






counting to infinity under the Split horizon with poisoned reverse

	距离	下一跳
B→D	2	E
C→D	2	E
E→D	无穷	

	距离	下一跳
B→D	无穷	
C→D	2	E
E→D	无穷	

	距离	下一跳
B→D	3	E
C→D	2	E
E→D	4	B

Unreachable message reached B but not reached C.





Triggered updates

- To get triggered updates, we simply add a rule that whenever a gateway changes the metric for a route, it is required to send update messages almost immediately, even if it is not yet time for one of the regular update message.



- RIP is a UDP-based protocol.
- Each host that uses RIP has a routing process that sends and receives datagrams on UDP port number 520.



作业

- 7, 8, 9, 10, 12, 16